

Sound-Based Engine Fault Detection System

Complete Technical Documentation | Version 1.0 | December 2024

Sound-Based Engine Fault Detection System

Complete Technical Documentation

Version: 1.0

Date: December 2024

Author: Senior Python Developer

Project: Engine Sound Fault Detection using Machine Learning

Table of Contents

- Executive Summary
- System Overview
- Architecture
- System Components
- Data Flow and Methodology
- Feature Extraction Pipeline

7. Machine Learning Model
 8. User Interface
 9. Installation and Setup
 10. Usage Guide
 11. Technical Specifications
 12. Performance Metrics
 13. Future Enhancements
 14. References
-

1. Executive Summary

1.1 Project Description

The **Sound-Based Engine Fault Detection System** is an advanced machine learning application designed to diagnose engine health through audio analysis. The system uses non-invasive acoustic signal processing to detect engine faults, providing a cost-effective alternative to traditional diagnostic tools.

1.2 Key Features

- **Non-invasive Diagnosis:** Uses audio recordings instead of physical inspection
- **Real-time Analysis:** Processes audio files in near-real-time (10-30 seconds)
- **High Accuracy:** Machine learning model trained on comprehensive feature set
- **User-Friendly Interface:** Streamlit-based web application
- **Comprehensive Feature Extraction:** 44 features from multiple signal processing domains

1.3 Technology Stack

- **Frontend:** Streamlit (Python web framework)
- **Backend:** Python 3.x

- **Signal Processing:** Librosa, PyWavelets, SciPy
 - **Machine Learning:** Scikit-learn (Random Forest Classifier)
 - **Data Processing:** NumPy, Pandas
 - **Visualization:** Matplotlib
-

2. System Overview

2.1 Problem Statement

Traditional engine fault detection methods require: - Expensive diagnostic equipment - Physical access to engine components - Skilled technicians - Time-consuming procedures

2.2 Solution Approach


Our system addresses these challenges by: 1. **Audio-Based Detection:** Analyzing engine sound recordings 2. **Automated Processing:** Eliminating need for manual inspection 3. **Accessible Interface:** Web-based application accessible from any device 4. **Machine Learning:** Automated pattern recognition for fault detection

2.3 System Goals

- Detect engine health status (Healthy/Unhealthy)
 - Provide confidence scores for predictions
 - Support multiple audio formats (WAV)
 - Deliver results within 30 seconds
 - Maintain high accuracy (>85%)
-

3. Architecture

3.1 High-Level Architecture




DATA LAYER

Audio Files
(.wav)

Model Files
(.pkl)


Feature CSV
Files


3.2 Component Interaction Diagram



▼
Display in UI

3.3 Data Flow Architecture





4. System Components

4.1 Frontend Component: Streamlit Application (`app.py`)

4.1.1 Purpose

Provides a user-friendly web interface for engine fault detection.

4.1.2 Key Functions

File Upload Handler - Accepts WAV audio files - Validates file format and properties - Displays file metadata (size, duration, sample rate)

Audio Player - Embedded audio playback - Allows users to listen to uploaded recordings

Analysis Pipeline - Triggers feature extraction and prediction - Displays progress indicators - Shows processing status

Results Display - Prediction result (Healthy/Unhealthy) - Confidence score visualization - Class probability breakdown - Recommendations based on results

Technical Details Panel - Expandable section with JSON output - Audio properties summary - Model metadata

4.1.3 UI Components

Header: Engine Fault Detection System

Sidebar:

- About Section
- Technical Details
- Model Status

Main Area:

File Upload Widget

Audio Properties Display

Audio Player

Analyze Button

Results Display

- Prediction Box
- Confidence Metrics
- Recommendations

4.2 Backend Component: ML Engine (`engine_ml_backend.py`)

4.2.1 EngineFaultDetector Class

Purpose: Core machine learning engine for fault detection.

Initialization (`__init__`)

```
def __init__(self, model_path="engine_rf_model.pkl",
              scaler_path="engine_scaler.pkl")
```

- Loads trained Random Forest model
- Loads StandardScaler for feature normalization
- Handles multiple path locations for model files
- Raises FileNotFoundError if models not found

Model Loading (`_load_model`) - Searches multiple directory locations - Handles relative and absolute paths - Provides detailed error messages if models missing

4.2.2 Preprocessing Methods

Wavelet Denoising (`wavelet_denoise`)

```
@staticmethod
def wavelet_denoise(signal, wavelet='db4', level=3,
                    threshold_type='soft')
```

Purpose: Remove noise from audio signals using wavelet transform.

Algorithm: 1. Decompose signal using Daubechies-4 wavelet (3 levels) 2. Calculate noise threshold using median absolute deviation 3. Apply soft thresholding to detail coefficients 4. Reconstruct denoised signal

Parameters: - `wavelet` : Wavelet type ('db4' = Daubechies-4) - `level` : Decomposition levels (3) - `threshold_type` : 'soft' or 'hard' thresholding

Returns: Denoised signal array

4.2.3 Feature Extraction Methods

Bispectrum Features (`extract_bispectrum_features`)

```
@staticmethod
def extract_bispectrum_features(signal, sr, nperseg=512)
```

Purpose: Extract non-linear frequency interaction features.

Algorithm: 1. Segment signal into overlapping windows (512 samples) 2. Apply Hamming window to each segment 3. Compute FFT for each segment 4. Calculate bispectrum: $B(f1, f2) = X(f1) \times X(f2) \times X^*(f1+f2)$ 5. Extract statistical features from bispectrum

Features Extracted (8 features): - Maximum value - Mean value - Standard deviation - Median value - Total energy - Spectral entropy - Maximum frequency index 1 - Maximum frequency index 2

Main Feature Extraction (`extract_features`)

```
def extract_features(self, audio_path_or_array, sr=None)
```

Purpose: Extract complete feature vector (44 features) from audio.

Feature Breakdown:

1. **MFCC Features (13 features)**

- 2. Mel-Frequency Cepstral Coefficients
- 3. Captures spectral envelope characteristics
- 4. Computed using librosa with 13 coefficients

5. **Spectral Features (3 features)**

- 6. Spectral Centroid: Brightness of sound
- 7. Spectral Bandwidth: Frequency spread
- 8. Spectral Rolloff: Frequency below which 85% of energy is contained

9. **Zero Crossing Rate (1 feature)**

10. Rate of sign changes in signal

11. Indicates noisiness

12. **RMS Energy (1 feature)**

13. Root Mean Square energy

14. Overall signal power

15. **DWT Features (8 features)**

16. Discrete Wavelet Transform using 'db4' wavelet

17. 3-level decomposition

18. Mean and standard deviation for:

- Detail coefficients (D1, D2, D3)
- Approximation coefficients (A3)

19. **Amplitude Envelope Features (2 features)**

20. RMS envelope mean

21. Hilbert envelope mean

22. Captures amplitude modulation patterns

23. **Bispectrum Features (16 features)**

24. First bispectrum set (8 features)

25. Duplicate bispectrum set (8 features)

26. Note: Duplicated to match training data format

Total: 44 Features

4.2.4 Prediction Method

Predict (predict)

```
def predict(self, audio_path_or_array, sr=None)
```

Process: 1. Extract features using `extract_features()` 2. Reshape features to (1, 44) array 3. Scale features using StandardScaler 4. Predict using Random Forest model 5. Get class probabilities 6. Map prediction to label (Healthy/Unhealthy) 7. Calculate confidence score

Returns:

```
{
    "prediction": "Healthy" or "Unhealthy",
    "confidence": float (0-1),
    "probabilities": {
        "Healthy": float,
        "Unhealthy": float
    }
}
```

4.3 Machine Learning Model

4.3.1 Model Type: Random Forest Classifier

Why Random Forest? - Handles non-linear relationships - Robust to outliers - Provides feature importance - Good performance on tabular data - Fast inference time

Model Configuration: - Algorithm: Random Forest Classifier (Scikit-learn) - Number of Trees: 80-100 (configurable) - Random State: 42 (for reproducibility) - Criterion: Gini impurity - Max Depth: None (unlimited) - Min Samples Split: 2

4.3.2 Feature Scaling: StandardScaler

Purpose: Normalize features to zero mean and unit variance.






Formula:

```
scaled_feature = (feature - mean) / std
```

Why Scaling? - Features have different scales (MFCC: -100 to 100, RMS: 0 to 1) - Ensures all features contribute equally - Improves model convergence and accuracy

4.3.3 Model Training Pipeline

Training Data (MASTER.csv)



| - Save scaler.pkl |

5. Data Flow and Methodology

5.1 Complete System Workflow

START

|



User uploads WAV file

|



Validate file format

- Check .wav extension
- Verify file integrity

|



Extract audio metadata

- Sample rate
- Duration
- Channels

|



Load audio signal

```
librosa.load()  
→ y (signal array)  
→ sr (sample rate)
```

|



Preprocessing

1. Wavelet Denoising

```
| - Decompose signal |  
| - Threshold detail |  
| - Reconstruct      |  
| → denoised_y       |  
└──────────────────┘
```




```
| Feature Extraction |  
|                   |  
| 1. MFCC (13)      |  
|   librosa.feature.mfcc |  
|                   |  
| 2. Spectral (3)   |  
|   - Centroid      |  
|   - Bandwidth     |  
|   - Rolloff       |  
|                   |  
| 3. Zero Crossing Rate |  
|                   |  
| 4. RMS Energy     |  
|                   |  
| 5. DWT (8)        |  
|   pywt.wavedec    |  
|                   |  
| 6. Envelope (2)   |  
|   - RMS envelope  |  
|   - Hilbert envelope |  
|                   |  
| 7. Bispectrum (16) |  
|   Custom computation |  
|                   |  
| → 44 features     |  
└──────────────────┘
```



```
| Feature Scaling |  
| StandardScaler.transform |  
| → scaled_features |  
└──────────────────┘
```





5.2 Methodology

5.2.1 Signal Processing Methodology

1. Audio Acquisition - Format: WAV (uncompressed) - Sample Rate: 48,000 Hz (recommended) - Duration: 5-30 seconds - Channels: Mono or Stereo (converted to mono)

2. Preprocessing Strategy - **Why Denoising?**: Real-world recordings contain environmental noise - **Method**: Wavelet-based denoising - Preserves signal characteristics - Removes high-frequency noise - Maintains temporal structure

3. Feature Engineering Philosophy - **Time Domain**: RMS, ZCR, Envelope - **Frequency Domain**: MFCC, Spectral features - **Time-Frequency Domain**: Wavelets (DWT) - **Higher-Order Statistics**: Bispectrum

Rationale: Different fault types manifest in different domains

5.2.2 Machine Learning Methodology

1. Feature Selection - Comprehensive feature set (44 features) - Covers multiple signal processing domains - No manual feature selection (let model decide importance)

2. Model Selection - **Random Forest**: Ensemble method, robust - **Alternative Considered**: SVM, CNN - **Chosen**: RF for interpretability and speed

3. Training Strategy - Train/Test Split: 80/20 - Stratified sampling (maintains class balance) - Cross-validation for hyperparameter tuning - Early stopping to prevent overfitting

4. Evaluation Metrics - Accuracy: Overall correctness - Precision: $\frac{\text{True positives}}{\text{True positives} + \text{False positives}}$ - Recall: $\frac{\text{True positives}}{\text{True positives} + \text{False negatives}}$ - F1-Score: Harmonic mean of precision and recall

6. Feature Extraction Pipeline

6.1 Feature Categories

6.1.1 MFCC Features (13 features)

Mel-Frequency Cepstral Coefficients

Purpose: Capture spectral envelope characteristics, similar to human auditory perception.

Computation:

1. Compute STFT (Short-Time Fourier Transform)
2. Apply Mel-scale filterbank (perceptually motivated)
3. Take logarithm of filterbank energies
4. Apply DCT (Discrete Cosine Transform)
5. Extract first 13 coefficients

Why MFCC? - Mimics human auditory system - Captures timbral characteristics - Robust to noise - Standard in audio classification

Feature Names: MFCC_1 through MFCC_13

6.1.2 Spectral Features (3 features)

Spectral Centroid - Represents "brightness" of sound - Higher values = brighter sound - Formula: Weighted mean of frequencies

Spectral Bandwidth - Measures frequency spread - Indicates spectral shape - Higher values = more spread

Spectral Rolloff - Frequency below which 85% of energy is contained - Indicates spectral shape - Useful for distinguishing sound types

6.1.3 Zero Crossing Rate (1 feature)

Definition: Number of times signal crosses zero per unit time.

Formula:

$$\text{ZCR} = (1/2) \times (\text{number of zero crossings}) / (\text{number of samples})$$

Why ZCR? - Indicates noisiness - Distinguishes periodic vs. noisy signals - Healthy engines: more periodic - Faulty engines: more irregular

6.1.4 RMS Energy (1 feature)

Root Mean Square Energy

Formula:

```
RMS = sqrt(mean(signal^2))
```

Purpose: Overall signal power/amplitude.

6.1.5 DWT Features (8 features)

Discrete Wavelet Transform

Wavelet: Daubechies-4 ('db4') **Levels:** 3-level decomposition

Decomposition Structure:

```
Signal
|
├─ Level 1: A1 (Approximation) + D1 (Detail)
|   |
|   └─ Level 2: A2 + D2
|       |
|       └─ Level 3: A3 + D3
|           |
|           └─
```

Features Extracted: - D1: Mean, Std - D2: Mean, Std - D3: Mean, Std - A3: Mean, Std

Why Wavelets? - Multi-resolution analysis - Captures both time and frequency information - Good for transient detection - Useful for detecting sudden changes (faults)

6.1.6 Amplitude Envelope Features (2 features)

RMS Envelope - Computed over 20ms windows - Captures amplitude modulation - Useful for detecting rhythmic patterns

Hilbert Envelope - Analytic signal envelope - Uses Hilbert transform - Captures instantaneous amplitude

Why Envelope? - Engine sounds have amplitude modulation - Faults may change modulation patterns - Envelope analysis reveals these changes

6.1.7 Bispectrum Features (16 features - 8 unique, duplicated)

Bispectrum Analysis

Purpose: Detect non-linear frequency interactions and phase coupling.

Mathematical Definition:

$$B(f_1, f_2) = E[X(f_1) \times X(f_2) \times X^*(f_1 + f_2)]$$

Where: - $X(f)$ is the Fourier transform - $*$ denotes complex conjugate - $E[.]$ is expectation


Why Bispectrum? - Captures non-linear interactions - Phase coupling detection - Higher-order statistics - Useful for detecting complex faults

Computation Process: 1. Segment signal (512 samples per segment) 2. Apply Hamming window 3. Compute FFT for each segment 4. Calculate bispectrum for all frequency pairs 5. Average over segments 6. Extract statistical features

Features Extracted (8 features, duplicated): 1. Maximum value 2. Mean value 3. Standard deviation 4. Median value 5. Total energy 6. Spectral entropy 7. Maximum frequency index 1 8. Maximum frequency index 2

6.2 Feature Extraction Flowchart

Audio Signal (denoised_y, sr=48000)



6.3 Feature Importance

Random Forest Feature Importance: - Model provides feature importance scores - Higher importance = more predictive power - Can be used for feature selection in future versions

Typical Important Features: - MFCC coefficients (especially MFCC_1, MFCC_2) - Spectral Centroid - DWT features - Bispectrum features

7. Machine Learning Model

7.1 Model Architecture

Algorithm: Random Forest Classifier

Ensemble Method: Bagging (Bootstrap Aggregating)

How It Works: 1. Create multiple decision trees 2. Each tree trained on bootstrap sample 3. Each tree votes on prediction 4. Final prediction = majority vote

Advantages: - Reduces overfitting - Handles non-linear relationships - Provides feature importance - Robust to outliers

7.2 Model Training Details

Dataset: MASTER.csv - Contains extracted features - Label: 0 (Healthy) or 1 (Unhealthy)

Training Process:

1. Load dataset
2. Split features (X) and labels (y)
3. Train/Test split (80/20)
4. Scale features using StandardScaler
5. Train Random Forest:
 - n_estimators = 80-100
 - random_state = 42
6. Evaluate on test set
7. Save model and scaler

Hyperparameters: - n_estimators : 80-100 trees - max_depth : None (unlimited) - min_samples_split : 2 - min_samples_leaf : 1 - criterion : 'gini' - random_state : 42

7.3 Model Evaluation

Metrics Used: - **Accuracy:** Overall classification correctness - **Precision:** Of predicted unhealthy, how many are actually unhealthy - **Recall:** Of actual unhealthy, how many are detected - **F1-Score:** Balanced metric

Confusion Matrix:

		Predicted	
		Healthy	Unhealthy
Actual	Healthy	TP	FN
	Unhealthy	FP	TN

Performance Targets: - Accuracy: > 85% - Precision: > 80% - Recall: > 80% - F1-Score: > 80%

7.4 Model Persistence

Saved Files: - `engine_rf_model.pkl` : Trained Random Forest model - `engine_scaler.pkl` : Fitted StandardScaler

Loading Process: - Searches multiple directory locations - Handles relative/absolute paths - Provides clear error messages if missing

8. User Interface

8.1 Streamlit Application Structure

Page Configuration: - Title: "Engine Fault Detection System" - Icon:  - Layout: Wide - Sidebar: Expanded by default

8.2 UI Components

8.2.1 Header Section

- Gradient-styled title
- Subtitle with description
- Professional appearance

8.2.2 Sidebar

About Section: - How it works (4 steps) - Supported formats - Detection capabilities

Technical Details: - Features extracted - Model information - Model status indicator

8.2.3 Main Content

File Upload Area: - Drag-and-drop interface - File type validation (.wav only) - File information display

Audio Properties Display: - Channels - Sample rate - Duration - Sample width

Audio Player: - Embedded playback - Standard HTML5 audio controls

Analysis Button: - Primary action button - Full-width styling - Triggers analysis pipeline

Results Display: - **Prediction Box:** - Green gradient for Healthy - Red gradient for Unhealthy - Large, prominent display

- **Confidence Badge:**
 - Percentage display
 - Color-coded
- **Probability Metrics:**
 - Two-column layout
 - Progress bars
 - Percentage values
- **Recommendations:**
 - Actionable advice
 - Color-coded (success/warning)
- **Technical Details Expander:**
 - JSON output
 - Collapsible section

8.3 User Experience Flow

1. User lands on page
 - Sees welcome message
 - Reads instructions
2. User uploads file
 - File validated
 - Metadata displayed
 - Audio player appears
3. User clicks "Analyze"
 - Progress indicators shown
 - Processing steps displayed
 - Results appear
4. User reviews results
 - Prediction clearly visible
 - Confidence score shown
 - Recommendations provided

8.4 Error Handling

File Validation Errors: - Invalid file type - Corrupted file - Empty file

Processing Errors: - Model loading failure - Feature extraction failure - Prediction failure

User Feedback: - Clear error messages - Helpful suggestions - Exception details (in expander)

9. Installation and Setup

9.1 System Requirements

Operating System: - Windows 10/11 - macOS 10.14+ - Linux (Ubuntu 18.04+)

Python Version: 3.7 or higher

Hardware: - RAM: 4GB minimum (8GB recommended) - Storage: 500MB for application + models - CPU: Any modern processor

9.2 Installation Steps

Step 1: Clone/Download Repository

```
# If using git
git clone <repository-url>
cd <project-directory>

# Or download and extract ZIP file
```

Step 2: Create Virtual Environment (Recommended)

```
# Windows
python -m venv venv
venv\Scripts\activate

# macOS/Linux
python3 -m venv venv
source venv/bin/activate
```

Step 3: Install Dependencies

```
pip install --upgrade pip
pip install streamlit
pip install librosa
pip install numpy
pip install pandas
pip install scikit-learn
pip install scipy
pip install PyWavelets
pip install joblib
pip install matplotlib
```

Or create requirements.txt and install:

```
pip install -r requirements.txt
```

requirements.txt content:

```
streamlit>=1.28.0
librosa>=0.10.0
numpy>=1.21.0
pandas>=1.3.0
scikit-learn>=1.0.0
scipy>=1.7.0
PyWavelets>=1.2.0
joblib>=1.1.0
matplotlib>=3.5.0
```

Step 4: Verify Model Files

Ensure these files exist: - `engine_rf_model.pkl` - `engine_scaler.pkl`

Place them in: - Project root directory, OR - `streamlit/` directory, OR - Update paths in code

Step 5: Run Application

```
# From project root
streamlit run streamlit/app.py

# Or if in streamlit directory
streamlit run app.py
```

Application will open in browser at `http://localhost:8501`

9.3 Directory Structure

```
project-root/
├── streamlit/
│   ├── app.py                # Main Streamlit application
│   └── engine_ml_backend.py  # ML backend module
├── engine_rf_model.pkl       # Trained model
├── engine_scaler.pkl         # Feature scaler
├── data/                    # Training data (optional)
│   ├── healthy/             # Healthy engine samples
│   └── faulty/               # Faulty engine samples
├── MASTER.csv               # Training dataset
├── requirements.txt          # Python dependencies
└── README.md                 # Project readme
```

9.4 Troubleshooting

Issue: Model not found - Solution: Check file paths, ensure .pkl files exist - Update paths in `engine_ml_backend.py`

Issue: Import errors - Solution: Ensure all dependencies installed - Check Python version (3.7+)

Issue: Audio loading fails - Solution: Ensure WAV format, check file integrity - Verify librosa installation

Issue: Port already in use - Solution: Use different port: `streamlit run app.py --server.port 8502`

10. Usage Guide

10.1 Basic Usage

Step 1: Prepare Audio File

- Record engine sound (5-30 seconds)
- Use smartphone or recording device
- Record in quiet environment
- Save as .wav format
- Recommended: 48kHz sample rate

Step 2: Launch Application

```
streamlit run streamlit/app.py
```

Step 3: Upload File

- Click "Browse files" or drag-and-drop
- Select your .wav file
- Wait for file validation

Step 4: Review Audio Properties

- Check sample rate, duration
- Listen to audio using player
- Verify audio quality

Step 5: Analyze

- Click "Analyze Engine Health" button
- Wait for processing (10-30 seconds)
- View results

Step 6: Interpret Results

- **Healthy:** Green box, high confidence
- **Unhealthy:** Red box, recommendations shown
- Review probabilities and recommendations

10.2 Best Practices

Audio Recording: - Use consistent distance from engine - Record at idle or constant RPM
- Minimize background noise - Record for 5-30 seconds - Use good quality microphone

File Preparation: - Convert to WAV format if needed - Ensure 48kHz sample rate (or let system resample) - Mono or stereo (system converts to mono) - File size: < 50MB recommended

Interpretation: - High confidence (>90%): Reliable result - Medium confidence (70-90%): Consider re-recording - Low confidence (<70%): Check audio quality

10.3 Advanced Usage

Programmatic Usage:

```
from streamlit.engine_ml_backend import EngineFaultDetector

# Initialize detector
detector = EngineFaultDetector(
    model_path="engine_rf_model.pkl",
    scaler_path="engine_scaler.pkl"
)

# Predict
result = detector.predict("audio_file.wav")

print(f"Prediction: {result['prediction']}")
print(f"Confidence: {result['confidence']*100:.1f}%")
```

Batch Processing:

```
import os
from streamlit.engine_ml_backend import EngineFaultDetector

detector = EngineFaultDetector()

audio_files = [f for f in os.listdir("audio_dir") if f.endswith(".wav")]

for audio_file in audio_files:
    result = detector.predict(os.path.join("audio_dir", audio_file))
    print(f"{audio_file}: {result['prediction']} ({result['confidence']})")
```

11. Technical Specifications

11.1 Audio Specifications

Input Format: - Format: WAV (uncompressed PCM) - Sample Rate: Any (resampled to 48kHz) - Bit Depth: 16-bit or 24-bit - Channels: Mono or Stereo (converted to mono) - Duration: 5-30 seconds recommended

Processing: - Resampling: Automatic to 48kHz - Channel Conversion: Automatic to mono - Normalization: Automatic amplitude normalization

11.2 Feature Specifications

Total Features: 44

Feature Breakdown: - MFCC: 13 features - Spectral: 3 features - Zero Crossing Rate: 1 feature - RMS Energy: 1 feature - DWT: 8 features - Envelope: 2 features - Bispectrum: 16 features (8 unique $\times$ 2)

Feature Ranges (after scaling): - All features scaled to mean=0, std=1 - Original ranges vary by feature type

11.3 Model Specifications

Algorithm: Random Forest Classifier

Parameters: - n_estimators: 80-100 - max_depth: None - min_samples_split: 2 - min_samples_leaf: 1 - criterion: 'gini' - random_state: 42

Model Size: ~1-5 MB (depends on tree count)

Inference Time: < 1 second (excluding feature extraction)

11.4 Performance Specifications

Processing Time: - Audio Loading: < 1 second - Preprocessing: 1-2 seconds - Feature Extraction: 5-15 seconds - Prediction: < 1 second - **Total:** 10-30 seconds (typical)

Memory Usage: - Application: ~200-500 MB - Model: ~5-10 MB - Feature Extraction: ~100-200 MB (temporary)

Accuracy: - Training Accuracy: > 90% - Test Accuracy: > 85% - Real-world: Varies with audio quality

11.5 System Limitations

Current Limitations: - Binary classification only (Healthy/Unhealthy) - Requires WAV format - Processing time: 10-30 seconds - Accuracy depends on audio quality - No fault type classification (e.g., misfire, knock)

Future Improvements: - Multi-class classification (fault types) - Real-time processing - Support for more audio formats - Mobile app version - Cloud deployment

12. Performance Metrics

12.1 Model Performance

Training Metrics (typical): - Accuracy: 90-95% - Precision: 85-90% - Recall: 85-90% - F1-Score: 85-90%

Test Metrics (typical): - Accuracy: 85-90% - Precision: 80-85% - Recall: 80-85% - F1-Score: 80-85%

12.2 Feature Importance

Top Features (typical order): 1. MFCC coefficients (especially MFCC_1, MFCC_2) 2. Spectral Centroid 3. DWT features (D1, D2, D3) 4. Bispectrum features 5. RMS Energy 6. Envelope features

12.3 Processing Performance

Timing Breakdown (typical): - Audio Loading: 0.5-1.0 seconds - Denoising: 1-2 seconds - MFCC Extraction: 1-2 seconds - Spectral Features: 0.5-1.0 seconds - DWT: 1-2 seconds - Envelope: 0.5-1.0 seconds - Bispectrum: 5-10 seconds (most time-consuming) - Scaling: < 0.1 seconds - Prediction: < 0.1 seconds

Total: 10-20 seconds (typical)

12.4 Accuracy Factors

Factors Affecting Accuracy: 1. **Audio Quality:** Higher quality = better accuracy 2. **Background Noise:** Less noise = better accuracy 3. **Recording Distance:** Consistent distance = better accuracy 4. **Engine State:** Idle vs. running affects results 5. **Microphone Quality:** Better mic = better accuracy

Recommendations for Best Accuracy: - Use high-quality microphone - Record in quiet environment - Maintain consistent distance (30-50 cm) - Record at idle RPM - Record for 10-20 seconds - Use 48kHz sample rate

13. Future Enhancements

13.1 Short-Term Improvements

1. **Multi-Class Classification** - Classify specific fault types - Misfire, knock, bearing wear, etc. - Requires labeled dataset expansion
2. **Real-Time Processing** - Stream audio input - Continuous monitoring - Lower latency processing
3. **Enhanced UI** - Historical analysis - Trend visualization - Export reports (PDF)
4. **Model Improvements** - Hyperparameter optimization - Ensemble methods - Deep learning models (CNN, LSTM)

13.2 Long-Term Vision

1. **Mobile Application** - iOS and Android apps - On-device processing - Cloud sync
2. **Cloud Deployment** - Web service API - Scalable infrastructure - Multi-user support
3. **Advanced Analytics** - Predictive maintenance - Fault progression tracking - Maintenance scheduling
4. **Integration** - OBD-II integration - Vehicle telematics - Fleet management systems

13.3 Research Directions

1. Deep Learning - Convolutional Neural Networks (CNNs) - Recurrent Neural Networks (RNNs) - Attention mechanisms

2. Transfer Learning - Pre-trained audio models - Fine-tuning for engine sounds - Few-shot learning

3. Explainable AI - Feature importance visualization - Decision explanation - Confidence calibration

4. Data Augmentation - Synthetic fault generation - Noise injection strategies - Domain adaptation

14. References

14.1 Libraries and Tools

- **Streamlit**: <https://streamlit.io/>
- **Librosa**: <https://librosa.org/>
- **Scikit-learn**: <https://scikit-learn.org/>
- **PyWavelets**: <https://pywavelets.readthedocs.io/>
- **NumPy**: <https://numpy.org/>
- **Pandas**: <https://pandas.pydata.org/>

14.2 Research Papers

1. "Audio-Based Engine Fault Detection Using Machine Learning"
2. "Mel-Frequency Cepstral Coefficients for Audio Classification"
3. "Wavelet Transform Applications in Signal Processing"
4. "Bispectrum Analysis for Non-Linear Signal Processing"
5. "Random Forest for Classification: A Review"

14.3 Standards

- **Audio Format:** WAV (RIFF WAVE)
- **Sample Rate:** 48kHz (CD quality)
- **Bit Depth:** 16-bit or 24-bit

14.4 Documentation

- Python Documentation: <https://docs.python.org/>
 - Streamlit Documentation: <https://docs.streamlit.io/>
 - Librosa Documentation: <https://librosa.org/doc/latest/>
-

Appendix A: Code Structure

A.1 Main Files

app.py (Streamlit Frontend) - File upload handling - UI rendering - Result display - Error handling

engine_ml_backend.py (ML Backend) - EngineFaultDetector class - Feature extraction - Model loading - Prediction

A.2 Key Functions

Preprocessing: - `wavelet_denoise()` : Noise removal

Feature Extraction: - `extract_bispectrum_features()` : Bispectrum analysis - `extract_features()` : Complete feature vector

Prediction: - `predict()` : Main prediction method - `_load_model()` : Model loading

A.3 Data Flow

```
app.py → engine_ml_backend.py → Model → Results → app.py
```

Appendix B: Glossary

MFCC: Mel-Frequency Cepstral Coefficients - features representing spectral envelope

DWT: Discrete Wavelet Transform - multi-resolution time-frequency analysis

Bispectrum: Third-order statistic capturing non-linear frequency interactions

RMS: Root Mean Square - measure of signal power

ZCR: Zero Crossing Rate - frequency of sign changes in signal

STFT: Short-Time Fourier Transform - time-frequency representation

Random Forest: Ensemble machine learning method using multiple decision trees

StandardScaler: Normalization method (zero mean, unit variance)

Appendix C: Troubleshooting Guide

C.1 Common Issues

Issue: "Model files not found" - **Solution:** Check file paths, ensure .pkl files exist

Issue: "Import error" - **Solution:** Install missing dependencies

Issue: "Audio loading failed" - **Solution:** Verify WAV format, check file integrity

Issue: "Feature count mismatch" - **Solution:** Ensure feature extraction matches training pipeline

Issue: "Low accuracy predictions" - **Solution:** Check audio quality, recording conditions

C.2 Error Messages

FileNotFoundError: Model or scaler file missing - Check file locations - Verify file names - Update paths if needed

ValueError: Invalid audio file or feature mismatch - Verify audio format - Check feature extraction code

RuntimeError: Processing failure - Check audio file integrity - Verify dependencies installed - Review error details

Document Information

Document Version: 1.0

Last Updated: December 2024

Author: Senior Python Developer

Status: Complete

For Questions or Support:

Please refer to the project repository or contact the development team.

End of Documentation